

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import pyplot as plt
from sklearn import preprocessing
from sklearn import metrics
from sklearn.preprocessing import MinMaxScaler
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn import svm
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier,
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report, f1_score, recall_score, precision_score
from mlxtend.plotting import plot_confusion_matrix
import lightgbm as lgb
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = pd.read_csv("Employee.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Education	JoiningYear	City	PaymentTier	Age	Gender	EverBenched	ExperienceInCurrentDomain
0	Bachelors	2017	Bangalore	3	34	Male	No	
1	Bachelors	2013	Pune	1	28	Female	No	
2	Bachelors	2014	New Delhi	3	38	Female	No	
3	Masters	2016	Bangalore	3	27	Male	No	
4	Masters	2017	Pune	3	24	Male	Yes	

## EDA

```
In [4]: df.rename(columns={'ExperienceInCurrentDomain': 'ECD'}, inplace=True)
```

In [5]: `df.describe()`

Out[5]:

	JoiningYear	PaymentTier	Age	ECD	LeaveOrNot
<b>count</b>	4653.000000	4653.000000	4653.000000	4653.000000	4653.000000
<b>mean</b>	2015.062970	2.698259	29.393295	2.905652	0.343864
<b>std</b>	1.863377	0.561435	4.826087	1.558240	0.475047
<b>min</b>	2012.000000	1.000000	22.000000	0.000000	0.000000
<b>25%</b>	2013.000000	3.000000	26.000000	2.000000	0.000000
<b>50%</b>	2015.000000	3.000000	28.000000	3.000000	0.000000
<b>75%</b>	2017.000000	3.000000	32.000000	4.000000	1.000000
<b>max</b>	2018.000000	3.000000	41.000000	7.000000	1.000000

In [6]: `df.shape`

Out[6]: (4653, 9)

In [7]: `df.columns`

Out[7]: Index(['Education', 'JoiningYear', 'City', 'PaymentTier', 'Age', 'Gender', 'EverBenched', 'ECD', 'LeaveOrNot'], dtype='object')

In [8]: `df.info()`


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4653 entries, 0 to 4652
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Education       4653 non-null   object
1   JoiningYear     4653 non-null   int64
2   City            4653 non-null   object
3   PaymentTier     4653 non-null   int64
4   Age             4653 non-null   int64
5   Gender          4653 non-null   object
6   EverBenched     4653 non-null   object
7   ECD             4653 non-null   int64
8   LeaveOrNot      4653 non-null   int64
dtypes: int64(5), object(4)
memory usage: 327.3+ KB
```

```
In [9]: df.isnull().sum()
```

```
Out[9]: Education      0
JoiningYear    0
City           0
PaymentTier    0
Age            0
Gender         0
EverBenched    0
ECD            0
LeaveOrNot      0
dtype: int64
```

```
In [10]: sns.heatmap(df.isnull())
```

```
Out[10]: <AxesSubplot:>
```




## Categorical Features

```
In [11]: labelencoder = preprocessing.LabelEncoder()
df['Education'] = labelencoder.fit_transform(df['Education'])
df['City'] = labelencoder.fit_transform(df['City'])
df['Gender'] = labelencoder.fit_transform(df['Gender'])
df['EverBenched'] = labelencoder.fit_transform(df['EverBenched'])
```

```
In [12]: plt.figure(figsize = (12,8))
sns.heatmap(df.corr() , annot=True,cmap=sns.dark_palette((250, 75, 60), input="h
```

Out[12]: <AxesSubplot:>




**Correlation is very weak between features !!!**

**Data visualization**

```
In [13]: sns.set_theme(palette="viridis")
fig, axs = plt.subplots(ncols=2,figsize=(16, 8))
(df['LeaveOrNot'].value_counts(normalize=True)*100).plot.pie(autopct='%1.1f%%', ax=axs[0])
(df['LeaveOrNot'].value_counts(normalize=True)*100).plot.bar(ax=axs[1])
```

Out[13]: <AxesSubplot:>




**0 >>> 65% staying**

**1 >>> 34% leaving**

**imbalnceing in the target**

```
In [14]: fig, axs = plt.subplots(ncols=2,figsize=(16, 8))
explode = [0.1,0.0,0.0,0.1]
labels = ["Male Stay", "Female stay", "Female Leave", "Male Leave"]
sns.histplot(data=df, x="Gender", hue="LeaveOrNot", multiple="dodge", shrink=6, ax=axs[0])
(df[['Gender', 'LeaveOrNot']].value_counts(normalize=True)*100).plot.pie(autopct='%1.1f%%', ax=axs[1])
```

Out[14]: <AxesSubplot:ylabel='None'>




**Male >>> 44% staying, 15% leaving**

**Female >>> 21% staying, 19% leaving**

```
In [15]: fig, axs= plt.subplots(figsize=(14, 8))  
sns.histplot(data=df, x="Age", hue="LeaveOrNot", multiple="dodge", shrink=.7)
```


```
Out[15]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```



**28 years is the highest and lowest in leaving and staying**

```
In [16]: fig, axs= plt.subplots(figsize=(14, 8))  
sns.histplot(data=df, x="ECD", hue="LeaveOrNot", multiple="dodge", shrink=0.7)
```

```
Out[16]: <AxesSubplot:xlabel='ECD', ylabel='Count'>
```




**2years exp. >>> highest in staying and leaving**




```
In [17]: sns.histplot(data=df, x="PaymentTier", hue="Education", multiple="dodge", shrink=
```

```
Out[17]: <AxesSubplot:xlabel='PaymentTier', ylabel='Count'>
```



```
In [18]: fig, axs = plt.subplots(ncols=2,figsize=(16, 8))
explode = [0.0,0.0,0.09,0.09,0.3,0.3]
labels = ["Bachelors, 0", "Bachelors, 1", "Masters, 0", "Masters, 1", "PHD, 0", "PHD, 1"]
sns.histplot(data=df, x="Education", hue="LeaveOrNot", multiple="dodge", shrink=0.4)
(df[['Education', 'LeaveOrNot']].value_counts(normalize=True)*100).plot.pie(autopct='%1.1f%%', explode=explode, labels=labels)
```

Out[18]: <AxesSubplot:ylabel='None'>




**Bachelors >>> 53% Staying , 24% leaving**

**Master >>> 9.2% leaving, 9.2% staying**

```
In [19]: fig, axs = plt.subplots(ncols=2,figsize=(16, 8))
explode = [0.1,0.1,0.0,0.0,0.0,0.0]
labels = ["Lowest Stay", "Lowest Leave", "Mid Leave", "Mid Stay", "High Stay", "High Leave"]
sns.histplot(data=df, x="PaymentTier", hue="LeaveOrNot", multiple="dodge", shrink=0.8)
(df[['PaymentTier', 'LeaveOrNot']].value_counts(normalize=True)*100).plot.pie(autopct='%1.1f%%', explode=explode, labels=labels)
```

Out[19]: <AxesSubplot:ylabel='None'>



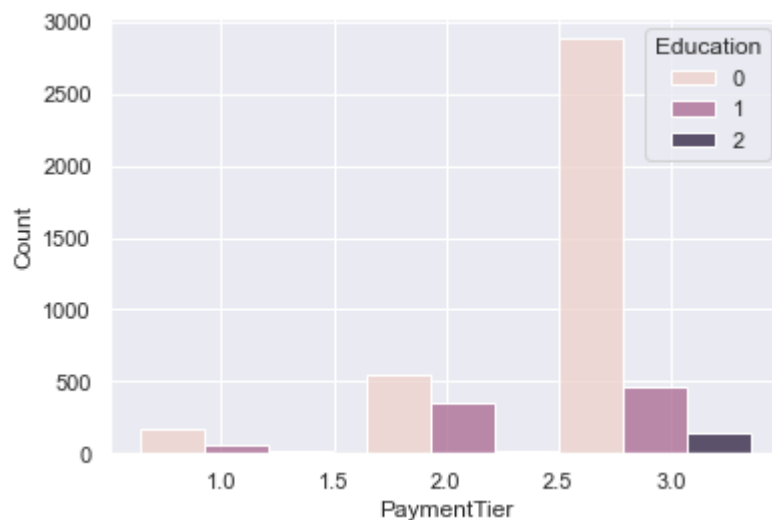
## The type of payment tier

-1: HIGHEST -2: MID LEVEL -3:LOWEST

tier payment >>> the lowest 54% staying ,20% leaving

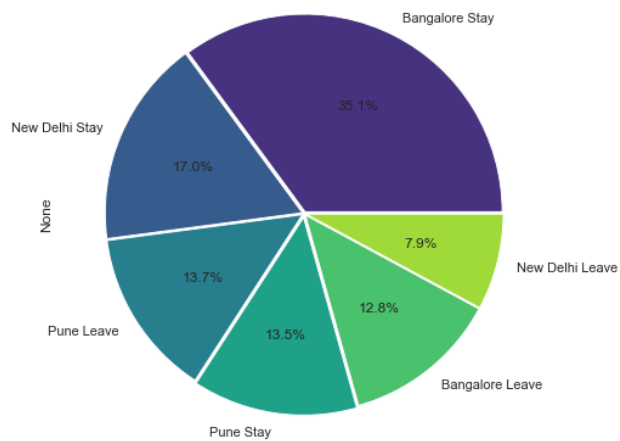

```
In [20]: sns.histplot(data=df, x="PaymentTier", hue="Education", multiple="dodge", shrink=0.8)
```

Out[20]: <AxesSubplot:xlabel='PaymentTier', ylabel='Count'>



```
In [21]: fig, axs = plt.subplots(ncols=2,figsize=(16, 8))
explode = [0.01,0.01,0.01,0.02,0.01,0.01]
labels = ["Bangalore Stay", "New Delhi Stay", "Pune Leave", "Pune Stay", "Bangalore Leave", "New Delhi Leave"]
sns.histplot(data=df, x="City", hue="LeaveOrNot", multiple="dodge", shrink=4, ax=axs[0])
(df[['City', 'LeaveOrNot']].value_counts(normalize=True)*100).plot.pie(autopct='%1.1f%%', ax=axs[1])
```

Out[21]: <AxesSubplot:ylabel='None'>




**Bangalore >>> 12.8% leaving**

**New Delhi >>> 8% leaving**


```
In [22]: sns.histplot(data=df, x="PaymentTier", hue="City", multiple="dodge", shrink=4)
```

Out[22]: <AxesSubplot:xlabel='PaymentTier', ylabel='Count'>




```
In [23]: fig, axs = plt.subplots(ncols=2, figsize=(16, 8))
explode = [0.09, 0.0, 0.0, 0.09, 0.09, 0.0, 0.0, 0.2, 0.0, 0.0, 0.0, 0.0, 0.5, .5]
sns.histplot(data=df, x="JoiningYear", hue="LeaveOrNot", multiple="dodge", shrink
(df[['JoiningYear', 'LeaveOrNot']].value_counts(normalize=True)*100).plot.pie(auto
```

Out[23]: <AxesSubplot:ylabel='None'>



```
In [25]: fig, axs = plt.subplots(ncols=2,figsize=(16, 8))
labels = ["No, 0", "No, 1", "Yes, 0", "Yes, 1"]
sns.histplot(data=df, x="EverBenched", hue="LeaveOrNot", multiple="dodge", shrink=0.8)
(df[['EverBenched', 'LeaveOrNot']].value_counts(normalize=True)*100).plot.pie(autopct='%1.1f%%', labels=labels)
```

Out[25]: <AxesSubplot:ylabel='None'>



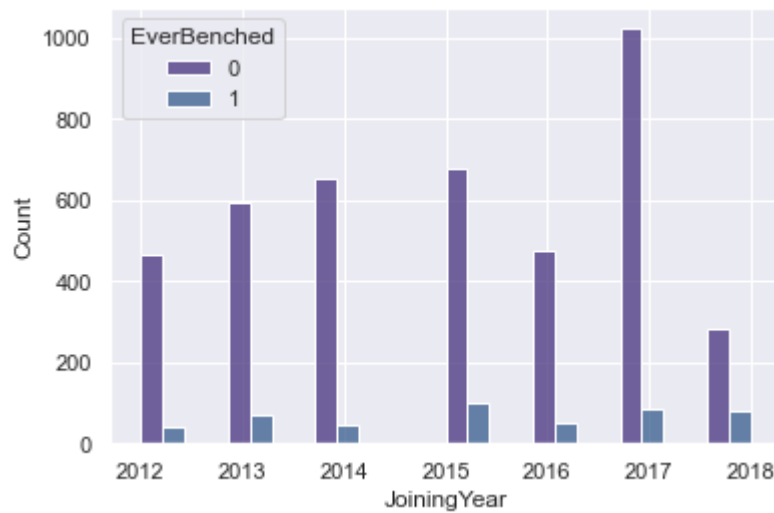
**it means ever kept out of the projects for 1 month or more**

**30% from not behanced leaving**

**5.6% behanced and not leaving**

```
In [26]: sns.histplot(data=df, x="JoiningYear", hue="EverBenched", multiple="dodge", shrink=1.2)
```

```
Out[26]: <AxesSubplot:xlabel='JoiningYear', ylabel='Count'>
```



## Normalization, Over sampling, and Train-Test-split

```
In [27]: sc = MinMaxScaler()
X = pd.DataFrame(sc.fit_transform(df.drop(["LeaveOrNot"], axis = 1)))
Y = df['LeaveOrNot'].values
```

```
In [28]: sm = SMOTE(random_state=42)
X,Y=sm.fit_resample(X, Y)
X_train,X_test,Y_train,Y_test = train_test_split(X,Y, test_size=0.2, random_state=42)
```

## Modeling

### 1-LGBMClassifier

```

In [29]: clf = lgb.LGBMClassifier()
clf.fit(X_train, Y_train)
Y_pred_test = clf.predict(X_test)
print("F1-Score:", metrics.f1_score(Y_test, Y_pred_test))
print("Accuracy:", metrics.accuracy_score(Y_test, Y_pred_test))
print("Precision:", metrics.precision_score(Y_test, Y_pred_test))
print("Recall:", metrics.recall_score(Y_test, Y_pred_test))
print("AUC:", metrics.roc_auc_score(Y_test, Y_pred_test))
print(classification_report(Y_test, Y_pred_test))
cutoff_grid = np.linspace(0.0, 1.0, 100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_test, pos_label=1)
confusion_matrix=confusion_matrix(Y_test,Y_pred_test)
plt.plot(FPR,TPR,c='red',linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plot_confusion_matrix(confusion_matrix,class_names=["not leaveing (0)","leaving(1)"])
plt.show()

```

F1-Score: 0.8771929824561403


Accuracy: 0.8854337152209493

Precision: 0.9328358208955224


Recall: 0.8278145695364238

AUC: 0.8847810711759789

	precision	recall	f1-score	support
0	0.85	0.94	0.89	618
1	0.93	0.83	0.88	604
accuracy			0.89	1222
macro avg	0.89	0.88	0.88	1222
weighted avg	0.89	0.89	0.89	1222








```
In [30]: def f_importances(coef, names, top=-1):
    imp = coef
    imp, names = zip(*sorted(list(zip(imp, names))))

    # Show all features
    if top == -1:
        top = len(names)

    plt.barh(range(top), imp[::-1][0:top], align='center')
    plt.yticks(range(top), names[::-1][0:top])
    plt.title('feature importances')
    plt.show()

features_names = ['Education', 'JoiningYear', 'City', 'PaymentTier', 'Age', 'Gender', 'EverBenched', 'ExperienceInCurrentDomain']
sns.set_theme(palette="Spectral")
f_importances(abs(clf.feature_importances_), features_names, top=8)
```



**in LGBMClassifier the best features for prediction is:**

**1- Age**

**2- Joining Year**

**3- Experience**

**4- City**

**5- Education**

**6- Payment Tier**

**7- Gender**

## 8- Benching

## 2-KNeighborsClassifier

```

In [31]: K=KNeighborsClassifier(n_neighbors=10)
K.fit(X_train, Y_train)
Y_pred_k = K.predict(X_test)
print("F1-Score:",metrics.f1_score(Y_test, Y_pred_k))
print("Accuracy:",metrics.accuracy_score(Y_test, Y_pred_k))
print("Precision:",metrics.precision_score(Y_test, Y_pred_k))
print("Recall:",metrics.recall_score(Y_test, Y_pred_k))
print("AUC:",metrics.roc_auc_score(Y_test, Y_pred_k))
print(classification_report(Y_test, Y_pred_k))
cutoff_grid = np.linspace(0.0,1.0,100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_k,pos_label=1)
plt.plot(FPR,TPR,c='red',linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plt.show()

```

F1-Score: 0.8141745894554884


Accuracy: 0.8240589198036007

Precision: 0.8517179023508138

Recall: 0.7798013245033113

AUC: 0.8235576201804583

	precision	recall	f1-score	support
0	0.80	0.87	0.83	618
1	0.85	0.78	0.81	604
accuracy			0.82	1222
macro avg	0.83	0.82	0.82	1222
weighted avg	0.83	0.82	0.82	1222



### 3-DecisionTreeClassifier

```

In [32]: D=DecisionTreeClassifier(max_depth=8,max_features=8,random_state=42)
D.fit(X_train, Y_train)
Y_pred_d = D.predict(X_test)
print("F1-Score:",metrics.f1_score(Y_test, Y_pred_d))
print("Accuracy:",metrics.accuracy_score(Y_test, Y_pred_d))
print("Precision:",metrics.precision_score(Y_test, Y_pred_d))
print("Recall:",metrics.recall_score(Y_test, Y_pred_d))
print("AUC:",metrics.roc_auc_score(Y_test, Y_pred_d))
print(classification_report(Y_test, Y_pred_d))
cutoff_grid = np.linspace(0.0,1.0,100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_test,pos_label=1)
plt.plot(FPR,TPR,c='red',linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plt.show()

```

F1-Score: 0.8313796212804329


Accuracy: 0.8469721767594108

Precision: 0.9128712871287129

Recall: 0.7632450331125827

AUC: 0.8460238110546733

	precision	recall	f1-score	support
0	0.80	0.93	0.86	618
1	0.91	0.76	0.83	604
accuracy			0.85	1222
macro avg	0.86	0.85	0.85	1222
weighted avg	0.86	0.85	0.85	1222




```
In [33]: def f_importances(coef, names, top=-1):
    imp = coef
    imp, names = zip(*sorted(list(zip(imp, names))))

    # Show all features
    if top == -1:
        top = len(names)

    plt.barh(range(top), imp[::-1][0:top], align='center')
    plt.yticks(range(top), names[::-1][0:top])
    plt.title('feature importances')
    plt.show()

features_names = ['Education', 'JoiningYear', 'City', 'PaymentTier', 'Age', 'Gender', 'EverBenched', 'ExperienceInCurrentDomain']
sns.set_theme(palette="coolwarm")
f_importances(abs(D.feature_importances_), features_names, top=8)
```



**in DecisionTreeClassifier the best features for prediction is:**

**1- Joining Year**

**2- Payment Tier**

**3- City**

**4- Education**

**5- Gender**

**6- Age**

**7- Experience**


**8- Benching**

**4-RandomForestClassifier**

```
In [34]: R=RandomForestClassifier(n_estimators=5,max_depth=8,max_features=8,random_state=42)
R.fit(X_train, Y_train)
Y_pred_r = R.predict(X_test)
print("F1-Score:",metrics.f1_score(Y_test, Y_pred_r))
print("Accuracy:",metrics.accuracy_score(Y_test, Y_pred_r))
print("Precision:",metrics.precision_score(Y_test, Y_pred_r))
print("Recall:",metrics.recall_score(Y_test, Y_pred_r))
print("AUC:",metrics.roc_auc_score(Y_test, Y_pred_r))
print(classification_report(Y_test, Y_pred_r))
cutoff_grid = np.linspace(0.0,1.0,100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_r,pos_label=1)
plt.plot(FPR,TPR,c='red',linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plt.show()
```

F1-Score: 0.8250226654578423  
Accuracy: 0.8420621931260229  
Precision: 0.9118236472945892  
Recall: 0.7533112582781457  
AUC: 0.8410569236374547

	precision	recall	f1-score	support
0	0.79	0.93	0.86	618
1	0.91	0.75	0.83	604
accuracy			0.84	1222
macro avg	0.85	0.84	0.84	1222
weighted avg	0.85	0.84	0.84	1222






```
In [35]: def f_importances(coef, names, top=-1):
    imp = coef
    imp, names = zip(*sorted(list(zip(imp, names))))

    # Show all features
    if top == -1:
        top = len(names)

    plt.barh(range(top), imp[::-1][0:top], align='center')
    plt.yticks(range(top), names[::-1][0:top])
    plt.title('feature importances')
    plt.show()

features_names = ['Education', 'JoiningYear', 'City', 'PaymentTier', 'Age', 'Gender', 'EverBenched', 'ExperienceInCurrentDomain']
sns.set_theme(palette="icefire")
f_importances(abs(R.feature_importances_), features_names, top=8)
```



**in RandomForestClassifier the best features for prediction is:**

**1- Joining Year**

**2- Payment Tier**

**3- City**

**4- Education**

**5- Age**

**6- Gender**

**7- Experience**

**8- Benching**

**5-GradientBoostingClassifier**

```
In [36]: G=GradientBoostingClassifier(n_estimators =5, max_depth =7, learning_rate = 0.3,
G.fit(X_train, Y_train)
Y_pred_g = G.predict(X_test)
print("F1-Score:",metrics.f1_score(Y_test, Y_pred_g))
print("Accuracy:",metrics.accuracy_score(Y_test, Y_pred_g))
print("Precision:",metrics.precision_score(Y_test, Y_pred_g))
print("Recall:",metrics.recall_score(Y_test, Y_pred_g))
print("AUC:",metrics.roc_auc_score(Y_test, Y_pred_g))
print(classification_report(Y_test, Y_pred_g))
cutoff_grid = np.linspace(0.0,1.0,100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_g,pos_label=1)
plt.plot(FPR,TPR,c='red',linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plt.show()
```

F1-Score: 0.8406840684068407


Accuracy: 0.8551554828150573

Precision: 0.9211045364891519

Recall: 0.7731788079470199

AUC: 0.8542269444265843

	precision	recall	f1-score	support
0	0.81	0.94	0.87	618
1	0.92	0.77	0.84	604
accuracy			0.86	1222
macro avg	0.86	0.85	0.85	1222
weighted avg	0.86	0.86	0.85	1222




```
In [37]: def f_importances(coef, names, top=-1):
    imp = coef
    imp, names = zip(*sorted(list(zip(imp, names))))

    # Show all features
    if top == -1:
        top = len(names)

    plt.barh(range(top), imp[::-1][0:top], align='center')
    plt.yticks(range(top), names[::-1][0:top])
    plt.title('feature importances')
    plt.show()

features_names = ['Education', 'JoiningYear', 'City', 'PaymentTier', 'Age', 'Gender', 'EverBenched', 'ExperienceInCurrentDomain']
sns.set_theme(palette="dark:salmon_r")
f_importances(abs(G.feature_importances_), features_names, top=8)
```



**in GradientBoostingClassifier the best features for prediction is:**

**1- Joining Year**

**2- City**

**3- Payment Tier**

**4- Education**

**5- Gender**

**6- Age**

**7- Experience**


**8- Benching**

**6-Support Vector Machine**

```
In [38]: rbf = svm.SVC(kernel='rbf', gamma=4,C=2).fit(X_train,Y_train)
rbf.fit(X_train, Y_train)
Y_pred_rbf = rbf.predict(X_test)
print("F1-Score:",metrics.f1_score(Y_test, Y_pred_rbf))
print("Accuracy:",metrics.accuracy_score(Y_test, Y_pred_rbf))
print("Precision:",metrics.precision_score(Y_test, Y_pred_rbf))
print("Recall:",metrics.recall_score(Y_test, Y_pred_rbf))
print("AUC:",metrics.roc_auc_score(Y_test, Y_pred_rbf))
print(classification_report(Y_test, Y_pred_rbf))
cutoff_grid = np.linspace(0.0,1.0,100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_rbf,pos_label=1)
plt.plot(FPR,TPR,c='red',linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plt.show()
```

F1-Score: 0.8224799286351473  
Accuracy: 0.837152209492635  
Precision: 0.8916827852998066  
Recall: 0.7632450331125827  
AUC: 0.8363150731905956

	precision	recall	f1-score	support
0	0.80	0.91	0.85	618
1	0.89	0.76	0.82	604
accuracy			0.84	1222
macro avg	0.84	0.84	0.84	1222
weighted avg	0.84	0.84	0.84	1222



7-ExtraTreesClassifier

```

In [39]: E = ExtraTreesClassifier(n_estimators=6,min_samples_split=10, random_state=0)
E.fit(X_train, Y_train)
Y_pred_e = E.predict(X_test)
print("F1-Score:",metrics.f1_score(Y_test, Y_pred_e))
print("Accuracy:",metrics.accuracy_score(Y_test, Y_pred_e))
print("Precision:",metrics.precision_score(Y_test, Y_pred_e))
print("Recall:",metrics.recall_score(Y_test, Y_pred_e))
print("AUC:",metrics.roc_auc_score(Y_test, Y_pred_e))
print(classification_report(Y_test, Y_pred_e))
cutoff_grid = np.linspace(0.0,1.0,100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_e,pos_label=1)
plt.plot(FPR,TPR,c='red',linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plt.show()

```

F1-Score: 0.858850174216028


Accuracy: 0.867430441898527

Precision: 0.90625

Recall: 0.8162251655629139

AUC: 0.8668504468591268

	precision	recall	f1-score	support
0	0.84	0.92	0.88	618
1	0.91	0.82	0.86	604
accuracy			0.87	1222
macro avg	0.87	0.87	0.87	1222
weighted avg	0.87	0.87	0.87	1222




```
In [40]: def f_importances(coef, names, top=-1):
    imp = coef
    imp, names = zip(*sorted(list(zip(imp, names))))

    # Show all features
    if top == -1:
        top = len(names)

    plt.barh(range(top), imp[::-1][0:top], align='center')
    plt.yticks(range(top), names[::-1][0:top])
    plt.title('feature importances')
    plt.show()

features_names = ['Education', 'JoiningYear', 'City', 'PaymentTier', 'Age', 'Gender',
                  'EverBenched', 'ExperienceInCurrentDomain']

f_importances(abs(E.feature_importances_), features_names, top=8)
```



**in ExtraTreesClassifier the best features for prediction is:**

**1- Joining Year**

**2- City**

**3- Payment Tier**

**4- Education**

**5- Gender**



**6- Age**

**7- Experience**


**8- Benching**

**8-LogisticRegression**

```
In [41]: L=LogisticRegression()
L.fit(X_train, Y_train)
Y_pred_1 = L.predict(X_test)
print("F1-Score:",metrics.f1_score(Y_test, Y_pred_1))
print("Accuracy:",metrics.accuracy_score(Y_test, Y_pred_1))
print("Precision:",metrics.precision_score(Y_test, Y_pred_1))
print("Recall:",metrics.recall_score(Y_test, Y_pred_1))
print("AUC:",metrics.roc_auc_score(Y_test, Y_pred_1))
print(classification_report(Y_test, Y_pred_1))
cutoff_grid = np.linspace(0.0,1.0,100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_1,pos_label=1)
plt.plot(FPR,TPR,c='red',linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plt.show()
```

F1-Score: 0.6448979591836734  
Accuracy: 0.644026186579378  
Precision: 0.6360708534621579  
Recall: 0.6539735099337748  
AUC: 0.644138858526758

	precision	recall	f1-score	support
0	0.65	0.63	0.64	618
1	0.64	0.65	0.64	604
accuracy			0.64	1222
macro avg	0.64	0.64	0.64	1222
weighted avg	0.64	0.64	0.64	1222



9-AdaBoostClassifier

```

In [42]: A=AdaBoostClassifier(n_estimators=5,learning_rate=.5)
A.fit(X_train, Y_train)
Y_pred_a = A.predict(X_test)
print("F1-Score:",metrics.f1_score(Y_test, Y_pred_a))
print("Accuracy:",metrics.accuracy_score(Y_test, Y_pred_a))
print("Precision:",metrics.precision_score(Y_test, Y_pred_a))
print("Recall:",metrics.recall_score(Y_test, Y_pred_a))
print("AUC:",metrics.roc_auc_score(Y_test, Y_pred_a))
print(classification_report(Y_test, Y_pred_a))
cutoff_grid = np.linspace(0.0,1.0,100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_a,pos_label=1)
plt.plot(FPR,TPR,c='red',linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plt.show()

```

F1-Score: 0.7368421052631579


Accuracy: 0.7667757774140753

Precision: 0.8329853862212944

Recall: 0.6605960264900662

AUC: 0.7655730941511819


	precision	recall	f1-score	support
0	0.72	0.87	0.79	618
1	0.83	0.66	0.74	604
accuracy			0.77	1222
macro avg	0.78	0.77	0.76	1222
weighted avg	0.78	0.77	0.76	1222



```
In [43]: def f_importances(coef, names, top=-1):
    imp = coef
    imp, names = zip(*sorted(list(zip(imp, names))))

    # Show all features
    if top == -1:
        top = len(names)

    plt.barh(range(top), imp[::-1][0:top], align='center')
    plt.yticks(range(top), names[::-1][0:top])
    plt.title('feature importances')
    plt.show()
    features_names = ['Education', 'JoiningYear', 'City', 'PaymentTier', 'Age', 'Gender',
                      'EverBenched', 'ExperienceInCurrentDomain']
    f_importances(abs(A.feature_importances_), features_names, top=8)
```



**in AdaBoostClassifier the best features for prediction is:**

**1- Joining Year**

**2- Payment Tier , Gender , City**

**10-GaussianNB**


```

In [44]: GNB = GaussianNB()
GNB.fit(X_train, Y_train)
Y_pred_gnb = GNB.predict(X_test)
print("F1-Score:", metrics.f1_score(Y_test, Y_pred_gnb))
print("Accuracy:", metrics.accuracy_score(Y_test, Y_pred_gnb))
print("Precision:", metrics.precision_score(Y_test, Y_pred_gnb))
print("Recall:", metrics.recall_score(Y_test, Y_pred_gnb))
print("AUC:", metrics.roc_auc_score(Y_test, Y_pred_gnb))
print(classification_report(Y_test, Y_pred_gnb))
cutoff_grid = np.linspace(0.0, 1.0, 100)
TPR = []
FPR = []
cutoff_grid
FPR, TPR, cutoffs = metrics.roc_curve(Y_test, Y_pred_gnb, pos_label=1)
plt.plot(FPR, TPR, c='red', linewidth=1.0)
plt.xlabel('False Positive')
plt.ylabel('True Positive')
plt.title('ROC Curve')
plt.show()

```

F1-Score: 0.6453781512605041  
 Accuracy: 0.6546644844517185  
 Precision: 0.6552901023890785  
 Recall: 0.6357615894039735  
 AUC: 0.6544503739900127

	precision	recall	f1-score	support
0	0.65	0.67	0.66	618
1	0.66	0.64	0.65	604
accuracy			0.65	1222
macro avg	0.65	0.65	0.65	1222
weighted avg	0.65	0.65	0.65	1222





```

In [45]: def testModel():
    edu = int(input("Please,Enter Education level\n0-Bachelors\n1-Masters\n2-PHD\n3-PhD\n"))
    year = int(input("Please,Enter Joining year\n"))
    city = int(input("Please,Enter City\n0-Bangalore\n1-New Delhi\n2-Pune\n3-Other\n"))
    payment = int(input("Please,Enter Payment Tier\n1-Highest\n2-Mid Level\n3-Lowest\n"))
    age = int(input("Please,Enter Age\n"))
    gender = int(input("Please,Enter Gender\n0-Female\n1-Male\n"))
    benched = int(input("Please,Enter if employee ever Benched\n0-No\n1-Yes\n"))
    exper = int(input("Please,Enter Experience In Current Domain\n"))
    data = [[edu,year,city,payment,age,gender,benched,exper]]
    data = sc.transform(data)
    pred = clf.predict(data)
    if (pred == 1):
        print("Employee Will Leave\n")
        sug = int(input("Would you like to know the suggestions to make him stay?\n"))
        if(sug == 1):
            from random import choice
            while (1):
                data[0][1] = choice(X[1])
                data[0][2] = choice(X[2])
                data[0][3] = choice(X[3])
                data[0][6] = choice(X[6])
                sugPred = clf.predict(data)
                if (sugPred == 0):
                    data = (sc.inverse_transform(data)).astype(int)
                    data = data.astype(str)
                    if (data[0][2] == "0"):
                        data[0][2] = "Bangalore"
                    elif (data[0][2] == "1"):
                        data[0][2] = "New Delhi"
                    elif (data[0][2] == "2"):
                        data[0][2] = "Pune"
                    elif (data[0][2] == "3"):
                        data[0][2] = "Other"
                    if (data[0][3] == "1"):
                        data[0][3] = "Highest"
                    elif (data[0][3] == "2"):
                        data[0][3] = "Mid"
                    elif (data[0][3] == "3"):
                        data[0][3] = "Lowest"
                    if (data[0][6] == "0"):
                        data[0][6] = "not"
                    elif (data[0][6] == "1"):
                        data[0][6] = ""
                    print("We suggest to deal with the employee like %s,\nmoving to %s with %s experience\n")
                    break
            if (pred == 0):
                print("Employee will not leave")
testModel()

```

```

Please,Enter Education level
0-Bachelors
1-Masters
2-PHD
3-PhD
0
Please,Enter Joining year
2018
Please,Enter City

```

```
0-Bangalore
1-New Delhi
2-Pune
0
Please,Enter Payment Tier
1-Highest
2-Mid Level
3-Lowest
3
Please,Enter Age
28
Please,Enter Gender
0-Female
1-Male
1
Please,Enter if employee ever Benched
0-No
1-Yes
1
Please,Enter Experience In Current Domain
2
Employee Will Leave

Would you like to know the suggetions to make him stay?
1-Yes
2-No
1
We suggest to deal with the employee like 2012,
move the employee to office in New Delhi,
make employee payment in Lowest level
and not make the employee benched
```

In [ ]: